

# QDP++ and Chroma

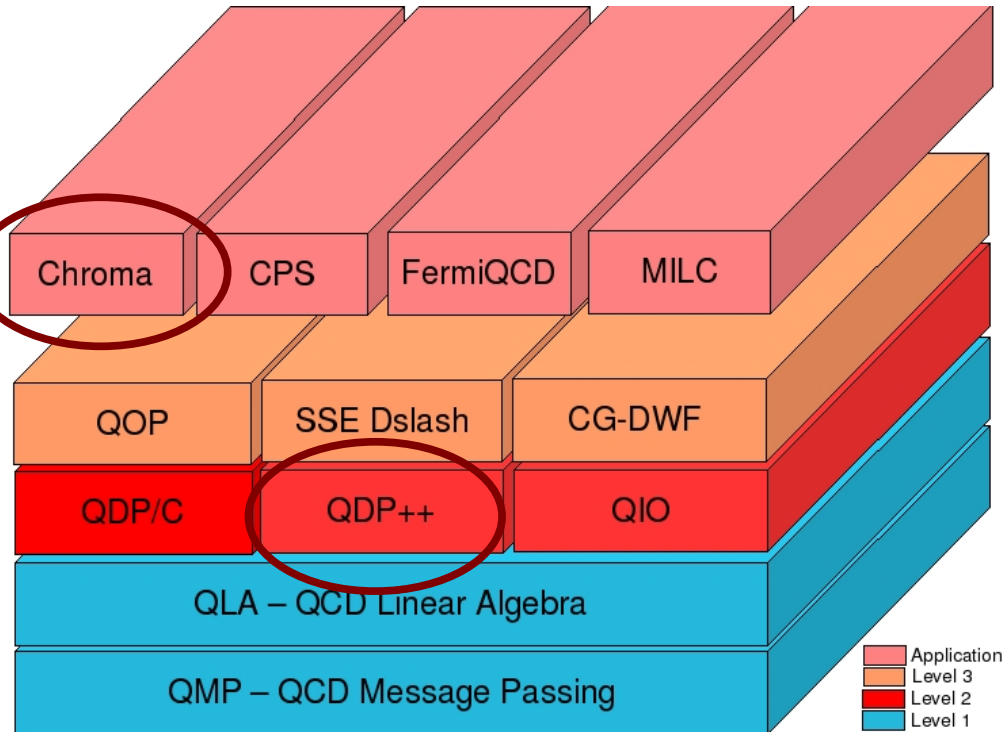
Balint Joo  
Jefferson Lab

HackLatt '08, Edinburgh

# Contents

- High Level View of QDP++ and Chroma
  - What they are, what they do etc
  - Software Maps
  - Getting, Building, Testing(?)
  - Using (?)
- QDP++ Details
  - Data Parallel Primitives
  - Subsets, Checkerboarding
  - I/O with XML, QIO

# QDP++ and Chroma What Are They



- Chroma is a lattice QCD framework written in C++
- Provides a library and some applications to do LQCD
- Relies on QDP++ to provide a data parallel platform
- Relies on 3<sup>rd</sup> party libraries for numerically intensive work

- QDP++ is platform for data parallel QCD computations
- It provides, expressions, shifts, memory management and I/O
- Relies on QIO library for binary I/O, QMP library for parallel communications

# Use of Chroma in the US

- USQCD Spectrum Project
  - Anisotropic Wilson/Clover Gauge Generation with (R)HMC algorithm
  - Wilson Clover Spectroscopy, photoproduction etc
- USQCD HASTE Project (now RBC-LHPC-UKQCD megacollaboration ?)
  - Domain Wall Fermion Forward/Sequential Propagator calculations
  - Hadron Structure (GPDA-s)
- USQCD NPLQCD Project
  - Domain Wall Fermion Propagator Calculations
- Algorithmic Studies
  - New inversion algorithms, preconditionings, dilution etc.

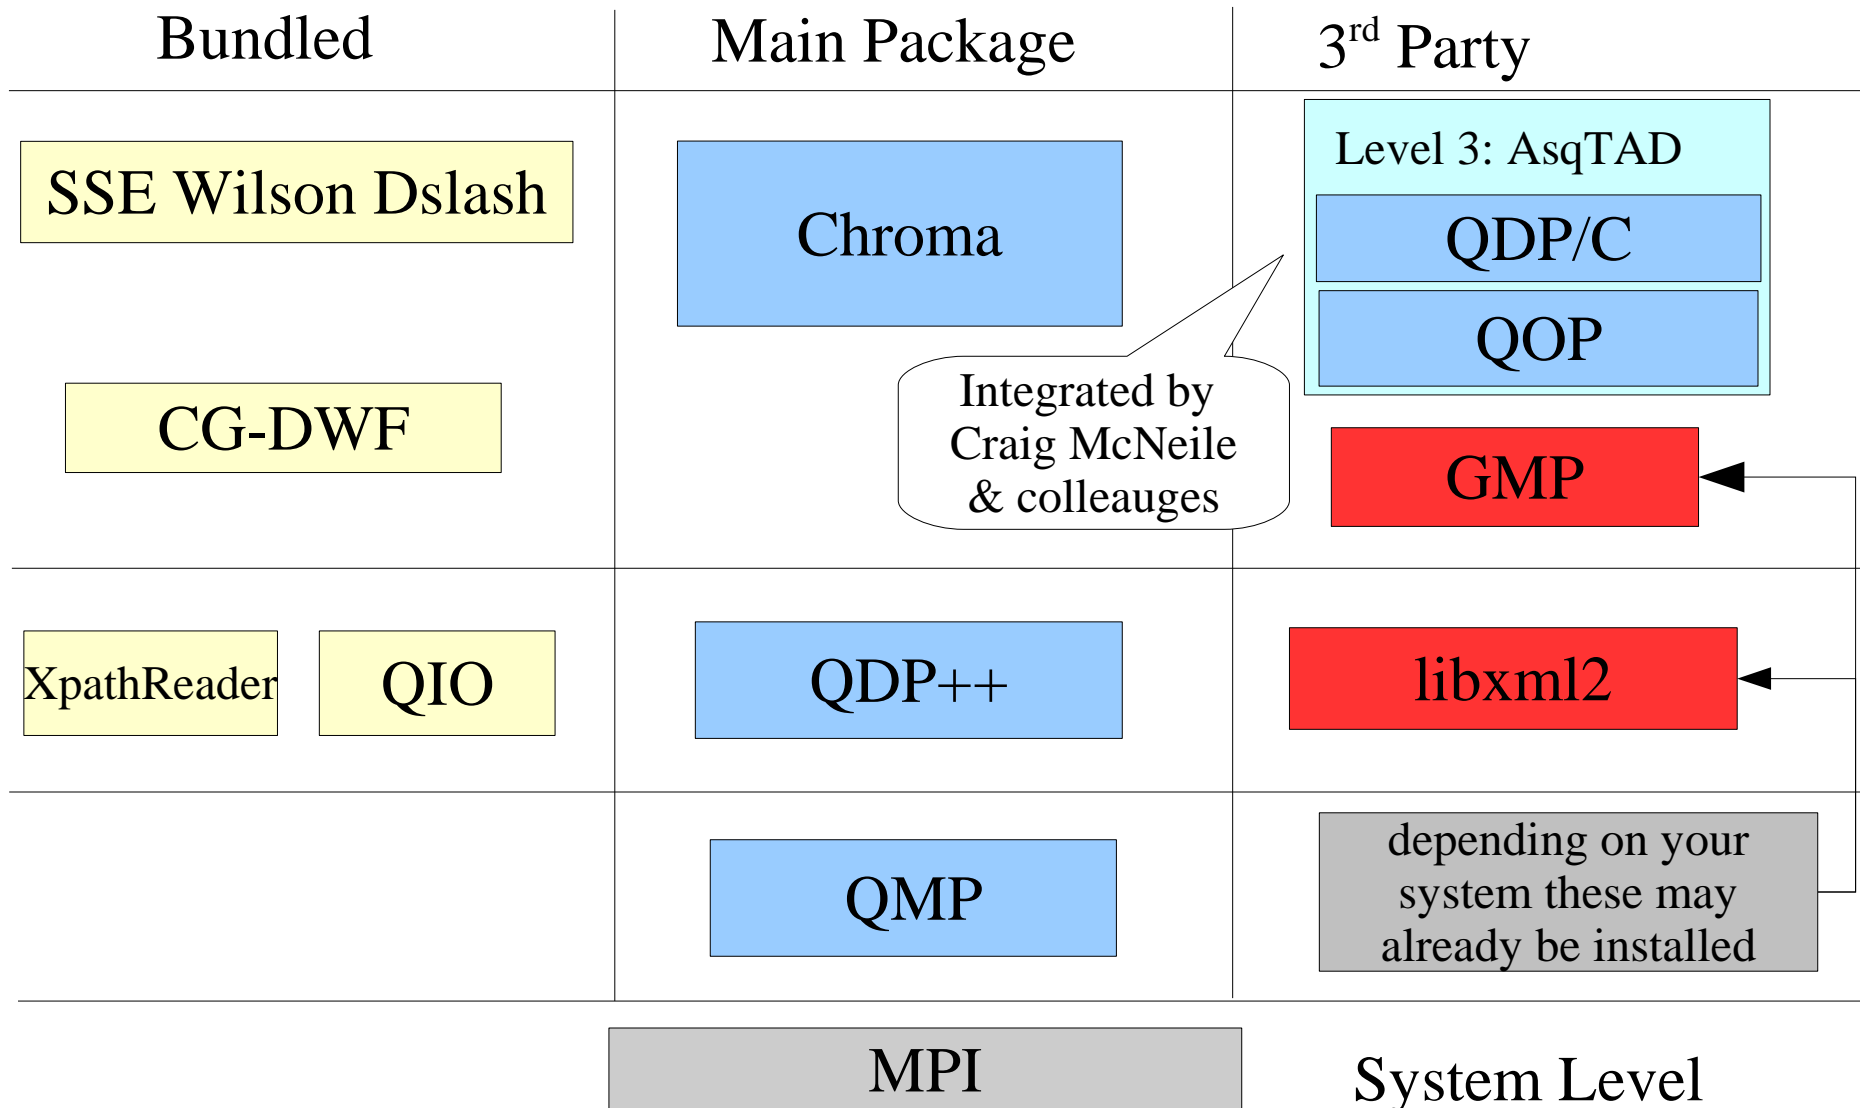
# Current List of Target Machines

- SSE2 Based Machines
  - Intel/Opteron workstations
  - Intel/Opteron clusters
  - Cray XT series supercomputers
  - QMP typically over MPI (MVIA-Mesh will deprecate)
  - Fast Dslash and some SSE optimized ops in QDP++
  - Fast DWF CG propagator solver from MIT
- BlueGene/L Machines
  - QMP typically over MPI (native QMP does exist)
  - Fast Bagel based Dslash from Peter Boyle
  - Fast QDP++ operations and Clover term coded using BAGEL Generator from Peter.
- QCDOC     Similar setup to BlueGene/L

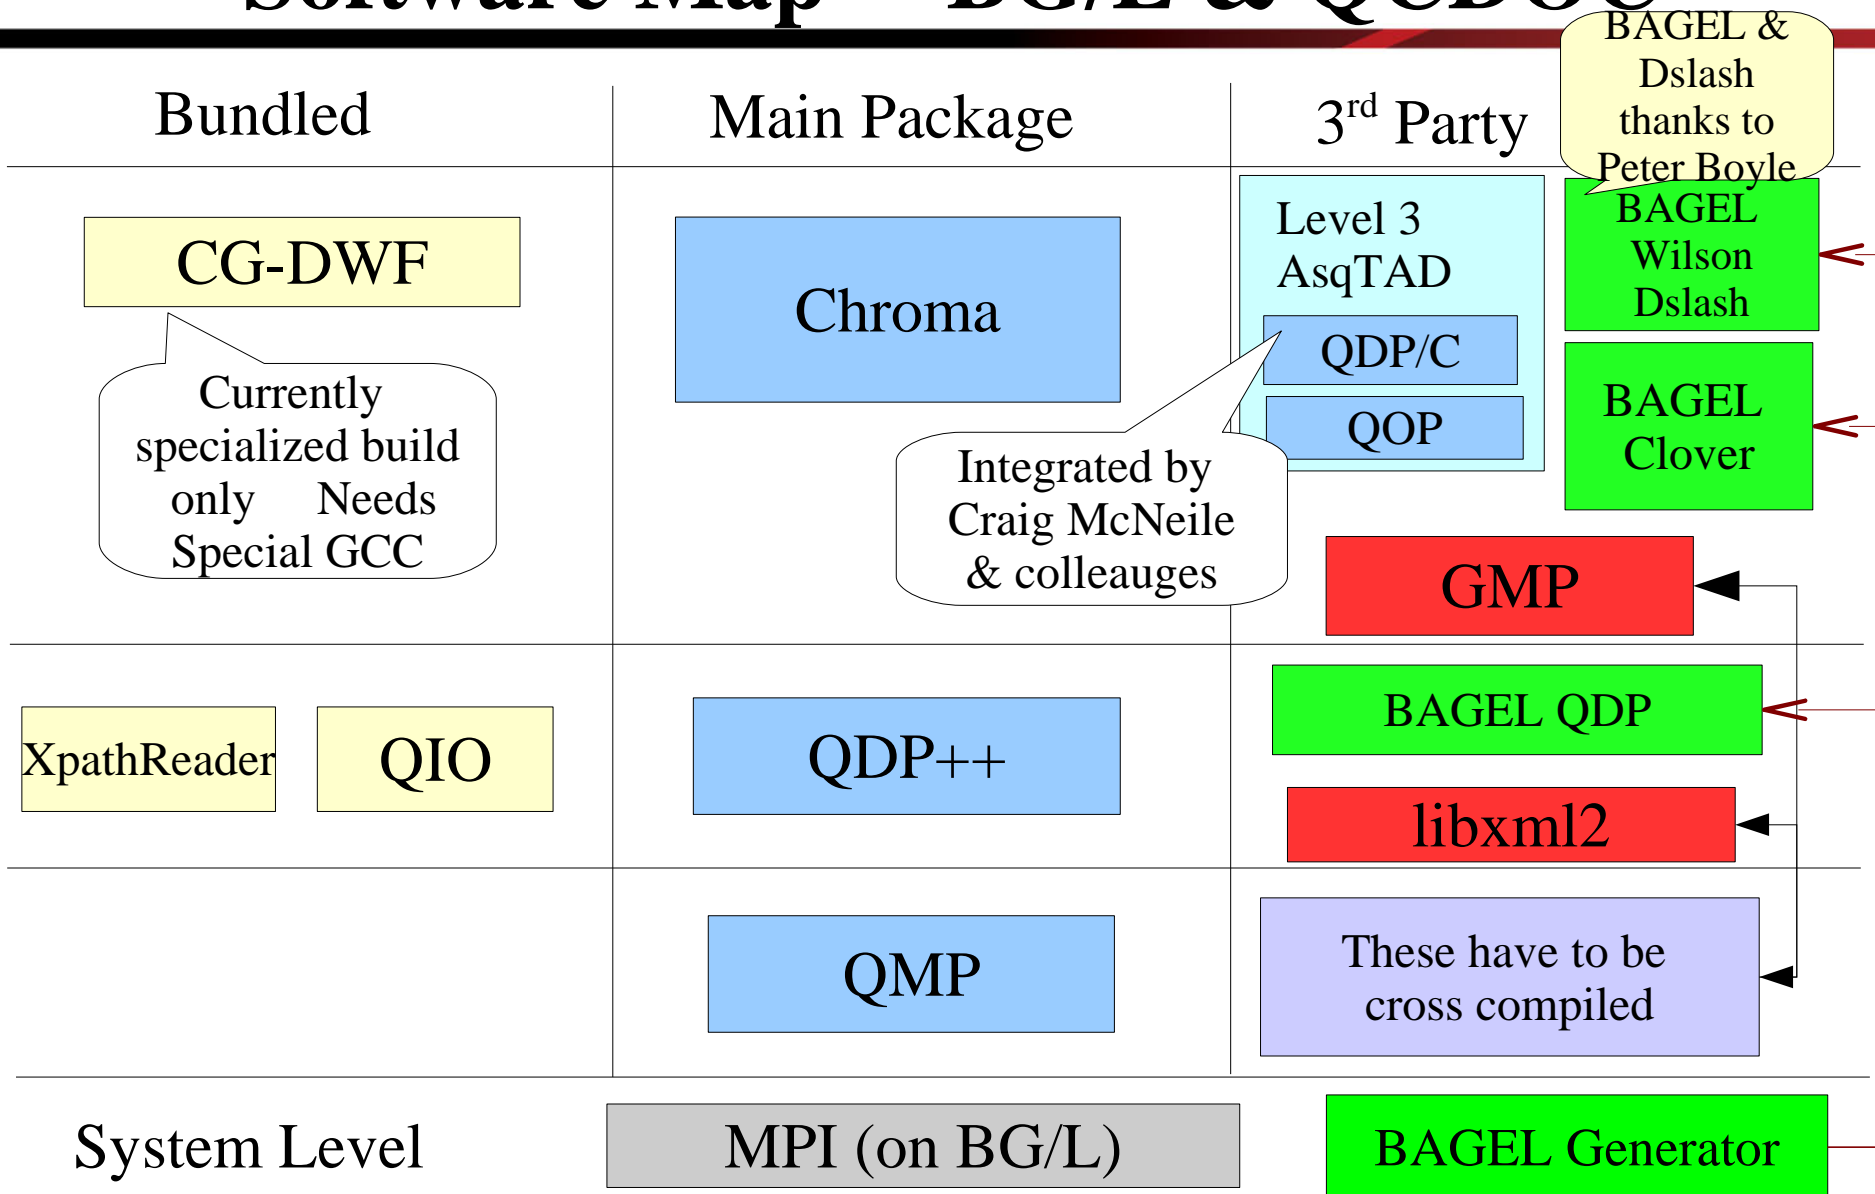
# Future/Development ports

- BlueGene/P
  - Start with BlueGene/L setup
  - Will do our best to integrate new components when/if we have access to them.
- Large Multi-Core Machines of the future
  - Ranger at TACC (port in progress)
  - whatever comes along?

# Software Map      SSE Based Systems



# Software Map BG/L & QCDOC





# Getting the bits and Pieces

- QMP, QDP++ and Chroma (and bundled packages) and BAGEL QDP, BAGEL Clover
  - from anonymous CVS:
    - Root :pserver:[anonymous@cvs.jlab.org](mailto:anonymous@cvs.jlab.org):/group/lattice/cvsroot
    - Modules: bagel\_qdp, bagel\_clover, qmp, qdp++, chroma
  - from USQCD Web Page:
    - <http://usqcd.jlab.org/usqcd-software>
- BAGEL and Wilson Dslash from Peter's web page
  - <http://www.ph.ed.ac.uk/~paboyle/bagel/Bagel.html>
- GMP from: <http://www.swox.com/gmp>
- LibXML2 from: <http://www.xmlsoft.org>

# Building

- Packages have a `configure` ; `make` ; `make install` style
  - but `configure` has many options that need to be set
- Simplest builds
  - use `mpicc`, `mpicxx`
  - build `qmp`, `qdp++`, `chroma` in sequence (`GMP`, `libxml` are already installed)
- Most complicated builds (`BlueGene/L` / `QCDOC`)
  - build `libxml2`, `GMP`, `bagel`, `bagel_qdp`, etc...
- This is a real pain to manage
- Jlab CVS module: `jlab-standard-chroma-build` attempts to automate the build workflow (download, configure, etc)
  - most build failures tend to be automake related :(
  - `QCDOC` poses some special problems...

# Compiler Versions etc

- We made the move to gcc-4.x series of compilers on SSE Based Platforms
- QMP and QDP++ should now also compile with Intel/Pathscale/PGI compilers    SSE may not be available with PGI    this needs investigation
- We've switched to automake-1.10 (automake versions tend to move along with our cluster upgrades)

# Some Configure Options for QDP++

- QDP++
  - --prefix=<install location>
  - --enable-parallel-arch=(scalar|parscalar)
  - --enable-precision=(single|double)
  - --with-libxml2=<location of libxml2 installation>
  - --with-qmp=<location of QMP installation>
- SSE Options
  - --enable-sse2
- BAGEL Related Options
  - --with-bagel-qdp=<location of Bagel QDP installation>
- QCDOC Specific Option
  - --enable-qcdoc (QCDOC Specific memory allocator)
- CXXFLAGS= -O2 -finline-limit=50000 CFLAGS= -O2

SZIN Leftovers:

scalar    workstation builds  
parscalar - parallel machine with scalar  
          processing elements - MPI/QMP

# Some Configure Options for Chroma

- Chroma
  - --prefix=<install location>
  - --with-qdp=<location of QDP++ installation>
  - --with-gmp=<location of GMP library>
- SSE Options
  - --enable-sse-wilson-dslash (SSE Only)
  - NEW {
    - --enable-cg-dwf=sse (SSE Domain Wall Solver)
    - --enable-cg-dwf-lowmem (Low memory DWF Solver)
    - --with-qmp=<location of QMP DWF Solver needs this>
- BAGEL related options
  - NEW {
    - --with-bagel-wilson-dslash=<location of BAGEL dslash>
    - --enable-bagel-wilson-dslash-sloppy <Sloppy Dslash>
    - --with-bagel-clover=<Location of BAGEL Clover>
- CXXFLAGS= CFLAGS= (disables default -g flag)

# Hints to Help You Build

- A (slightly out of date) set of installation notes on [the USQCD Web Site](#)
- The jlab-standard-chroma-build package in the Chroma CVS server
  - `cvs checkout jlab-standard-chroma-build`
  - `cd jlab-standard-chroma-build`
  - Need to configure this with `./configure`
  - Useful Options are
    - `--enable-install-root=<root of install tree>`
    - `--enable-qdp-version=<version of QDP++ to use>`
    - `--enable-chroma-version=<version of Chroma to use>`
    - `--enable-parallel-make=<no of threads for make>`

# More on Jlab Standard Chroma Build

- The standard build is organised in directories as:
  - package/machine/build-type
  - eg: Single precision, parscalar build of QDP++ for Jlab 7n machine: qdp++/ib-7n/parscalar-ib-7n
  - eg: Scalar build of Chroma: chroma/scalar/scalar
  - These triplets map directly onto directory names
    - `cd qdp++/ib-7n/parscalar-ib-7n`
    - There is a build/ subdirectory here.
      - `cd ./build/`
      - There is `configure.sh` script here.
        - In the `configure.sh` script are useful flags for this target
  - You can use this to get a handle on what flags you want where

# More on Jlab Standard Chroma Build

- Can actually build the code with this too
  - look at eg: `xt3_build.sh` complete the configure line then run on your favourite Cray XT machine
  - look at eg: `bgl_build.sh` complete the configure line then run on your favourite BG/L
- Logs of builds kept in `./logs/package`
- Successful builds installed in
  - `<install-root>/package/version/build-type`
  - eg: `/home/bjoo/install/qdp++/qdp1-25-1/parscalar-ib-7n`
- `./build.sh package/machine/build-type` will attempt to download package source for you
- Package source is downloaded to `package/package` subdirectory of `jlab-standard-chroma-build`
  - In case you need the source to run automake etc.



# Testing

- Chroma contains many regression tests (over 100) mostly in
  - `chroma/tests/chroma` (main chroma measurement tests)
  - `chroma/tests/t_leapfrog` (MD tests)
  - `chroma/tests/hmc` (HMC tests)
- Test comprised of:
  - executable (`chroma`, `hmc`, `t_leapfrog` etc)
  - input param file (XML)
  - expected output file (XML)
  - output comparison metric file (XML)
    - used to check code output against expected output
- Need `xmldiff` utility to check output against expected output
- Need some infrastructure to run through all tests

# How to run the tests

- Have xmldiff installed and on your PATH before configuring chroma
- After chroma is configured it will produce a script called RUN you may need to edit this. In the case of a scalar machine it may be as simple as
  - `#!/bin/bash`
  - `$*`
- For an MPI machine with 4 CPUs you may want it to look like
  - `#!/bin/bash`
  - `mpirun -np 4 $*`
- Then in an interactive session you should run 'make xcheck'
- There are complicated ways to run this through batch queues as well...

# Linking against an installed Chroma

- Suppose chroma is installed in `/foo/chroma`
- Use script `chroma-config` in `/foo/chroma/bin`
  - `CXX=`chroma-config --cxx``
  - `CXXFLAGS=`chroma-config --cxxflags``
  - `LDFLAGS=`chroma-config --ldflags``
  - `LIBS=`chroma-config --libs``
- Compile your program (`prog.cc`) with:
  - `$(CXX) $(CXXFLAGS) prog.cc $(LDFLAGS) $(LIBS)`
  - NB: Ordering of flags may be important.

# Running Chroma Applications

- Measurement Application: **chroma**
- Gauge Generation Applications: **hmc** and **purgaug**
- Installed in same place as chroma-config
  - eg: /foo/chroma/bin
- Typical usage flags (-i, -o, -geom):
  - **./chroma -i in.xml -o out.xml -geom Px Py Pz Pt**
  - **in.xml**     Input Parameter XML File
  - **out.xml**     Output XML Log File
  - **Px Py Pz Pt** the (possibly virtual) Processor Geometry  
(eg -geom 4 4 8 8 for QCDOC Rack)

# Chroma XML Input File Structure

```
<?xml version= 1.0 encoding= UTF-8 ?>
<chroma>
<annotation>Your annotation here</annotation>
<Param>
```

Array of Measurements (Tasks)

Task (array element)

Task Name

Task specific  
parameters

Named Objects  
(communicate between tasks  
-- like in memory files)

Global Lattice Size

Input Configuration to use as  
default\_gauge\_field

Task output  
XML file

```
<InlineMeasurements>
  <elem>
    <Name>MAKE_SOURCE</Name>
    <Frequency>1</Frequency>
    <Param/>
    <NamedObject>
      <gauge_id>default_gauge_field</gauge_id>
      <source_id>sh_source_0</source_id>
    </NamedObject>
  </elem>
  <elem>
    <Name>PROPAGATOR</Name>
    <Frequency>1</Frequency>
    <Param/>
    <NamedObject>
      <gauge_id>default_gauge_field<gauge_id>
      <source_id>sh_source_0</source_id>
      <prop_id>sh_prop_0</prop_id>
    </NamedObject>
    <xml_file>./prop_out.xml<xml_file>
  </elem>
</InlineMeasurements>
<nrow>4 4 4 8</nrow>
```

```
</Param>
```

```
<RNG/>
```

```
<Cfg>
  <cfg_type>SCIDAC</cfg_type>
  <cfg_file>foo.lime</cfg_file>
</Cfg>
```

```
</chroma>
```

# Where to find XML examples

- Measurement Tasks for:
- Numerous Measurement Task Examples in
  - chroma/tests/chroma/hadron/
  - sources, smearings, propagators, spectroscopy, 3pt functions, eigenvalues
  - Reading and Writing Named Objects
- Also MD and HMC input files in
  - chroma/tests/t\_leapfrog
  - chroma/tests/hmc
- Input file names usually contain the string ini

# QDP++ Details

- QDP++ Is the Foundation of Chroma
- It provides
  - way to write the maths of lattice QCD without looping over indices (expressions)
  - Custom memory allocation possibilities
  - I/O facilities (XML and Binary)
- You can do Lattice QCD just in QDP++ without Chroma
  - [See: Lectures from the 2007 INT Lattice Summer School](#)
  - but you'd need to write a whole lot of infrastructure that comes for free with Chroma

# QDP++ Templated Types

- In QDP++ tries to capture the index structure of our lattice fields:

	Lattice	Spin	Colour	Reality	BaseType
Real	Scalar	Scalar	Scalar	Real	REAL
LatticeColorMatrix	Lattice	Scalar	Matrix(Nc,Nc)	Complex	REAL
LatticePropagator	Lattice	Matrix(Ns,Ns)	Matrix(Nc,Nc)	Complex	REAL
LatticeFermionF	Lattice	Vector(Ns)	Vector(Nc)	Complex	REAL32
DComplex	Scalar	Scalar	Scalar	Complex	REAL64

- To do this we use C++ templated types:

```
typedef OScalar < PScalar < PScalar< RScalar <REAL> > > > Real;
typedef OLattice< PScalar < PColorMatrix< RComplex<REAL>, Nc> > > LatticeColorMatrix;
typedef OLattice< PSpinMatrix< PColorMatrix< RComplex<REAL>, Nc>, Ns> > LatticePropagator;
```

- QDP++ and Portable Expression Template Engine:
  - Provide expressions and recursion through type structure



# QDP++ and Expressions

" Use convenient mathematical expressions:

```
LatticeFermion x,y,z;
```

```
Real a = Real(1);
```

```
gaussian(x);
```

```
gaussian(y);
```

```
z = a*x + y;
```

```
int mu, nu;
```

```
multi1d<LatticeColorMatrix> u(Nd);
```

```
Double tmp = sum( real( trace( u[mu]
```

```
*shift(u[nu], FORWARD, mu)
```

```
* adj( shift(u[mu], FORWARD, nu) )
```

```
* adj(u[nu])
```

```
)
```

```
)
```

```
);
```

Wowee! No indices or for loops!

multi1d<T> =  
1 dimensional array  
of T

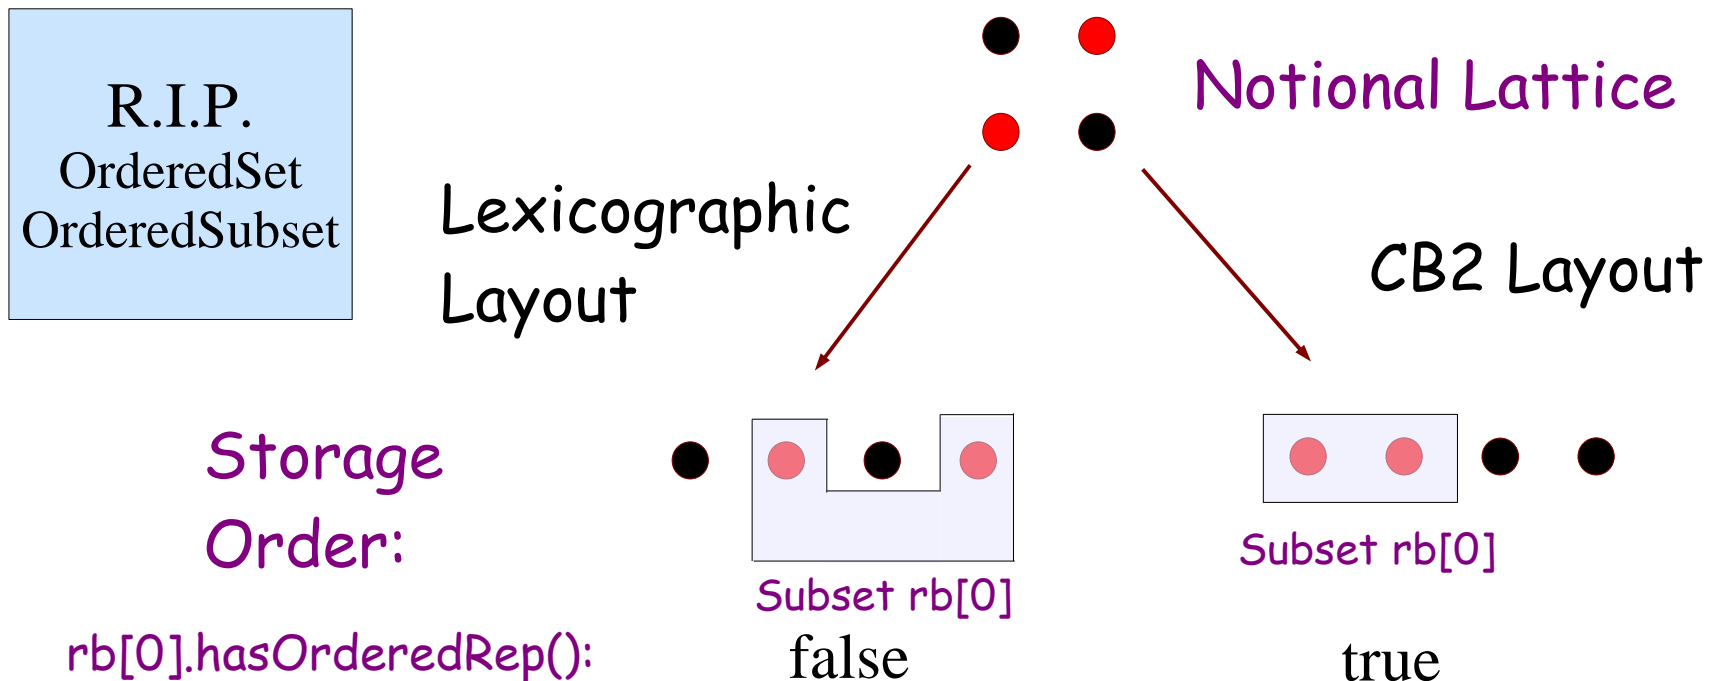
Lattice Wide Shifts  
FORWARD =  
from x+mu

Functions

but array indices are explicit

# Subsets and Layouts

- Subset: Object that identifies a subset of sites
- Can be predefined: eg rb is red-black colouring
- Can be contiguous or not (`s.hasOrderedRep()` == true or not)
- Layout is an ordering of sites in memory (compile time choice)
- Same subset may be contiguous in one layout and not in another



# Using Subsets

In expressions, subset index is on the target

```
bar[rb[1]] = foo; // (Copy foo's rb[1] subset to bar's)
```

- Users can define new sets
- Layout chosen at compile time with configure switch
  - Default is CB2 (2 colour checkerboard) layout
- Layout needs to be initialized on entry to QDP++

```
multi1d<int> nrow(4);  
nrow[0]=nrow[1]=nrow[2]=4; nrow[3]=8;  
Layout::setLattSize(nrow);  
Layout::create();
```

# QDP++ and XML

- No static data binding in QDP++
- Treat documents as amorphous (contain anything)
- Interrogate documents using XPath.
- Root of Path expression is context node

root node → `<?xml version= 1.0 encoding= UTF-8  
<foo>  
    <bar>  
        <fred>6</fred>  
        <jim>7 8 9</jim>  
    </bar>  
</foo>`

From root: `/foo/bar/fred`

from /bar: `./fred`

→ `<fred>6</fred>`  
→ `<jim>7 8 9</jim>`

`</bar>`  
`</foo>`

# Reading XML

```
XMLReader r( filename );
```

Array markup of  
non-simple types

```
Double y;
```

```
multild<Int> int_array;
```

```
multild<Complex> cmp_array;
```

```
try {
```

```
    read(r, /foo/cmp_array , cmp_array);
```

Absolute Path Read

New Context

```
    XMLReader new_r(r, /foo/bar );
```

Relative Paths

```
    read(new_r, ./int_array , int_array);
```

```
    read(new_r, ./double , y);
```

```
}
```

```
catch( const std::string& e) {
```

```
    QDPIO::cerr << Caught exception:
```

```
    << e <<endl;
```

```
    QDP_abort(1);
```

```
}
```

QDP++ error  
"stream"

```
<?xml version= 1.0  
      encoding= UTF-8 ?>
```

```
<foo>
```

```
  <cmp_array>
```

```
    <elem>
```

```
      <re>1</re>
```

```
      <im>-2.0</im>
```

```
    </elem>
```

```
  <elem>
```

```
    <re>2</re>
```

```
    <im>3</im>
```

```
  </elem>
```

```
</cmp_array>
```

```
<bar>
```

```
  <int_array>2 3 4 5</int_array>
```

```
  <double>1.0e-7</double>
```

```
</bar>
```

```
</foo>
```

Array markup  
for simple types

# Writing XML

```
// Write to file
XMLFileWriter foo( ./out.xml );
push(out, rootTag );
int x=5;
Real y=Real(2.0e-7);
write(foo, xTag , x);
write(foo, yTag , y);
pop(out);
```

<?xml version= 1.0 ?>  
<rootTag>  
<xTag>5</xTag>  
<yTag>2.0e-7</yTag>  
</rootTag>

```
// Write to Buffer
XMLBufferWriter foo_buf;
push(foo_buf, rootTag );
int x = 5;
Real y = Real(2.0e-7);
write(foo_buf, xTag , x);
write(foo_buf, yTag , y);
pop(foo_buf);
QDPIO::cout << Buffer contains <foo_buf.str()
               << endl;
```

get buffer content

<foo\_buf.str()

# QIO and LIME Files

Private File XML Data
User File XML Data
Private Record XML Data
User Record XML Data
Record Binary Data
Checksum Record
Private Record XML Data
User Record XML Data
Record Binary Data
Checksum Record
Private Record XML Data
User Record XML Data
Record Binary Data
Checksum Record

HEADER



Message 1  
Record 1

Message 1  
Record 2

Message 2  
Record 1

- QIO works with record oriented LIME files
- LIME files made up of messages
- messages are composed of
  - File XML records
  - Record XML records
  - Record Binary data
- SciDAC mandates checksum records
- ILDG mandates certain records


# Three modes of writing in QIO



● I/O process

● non I/O process

file

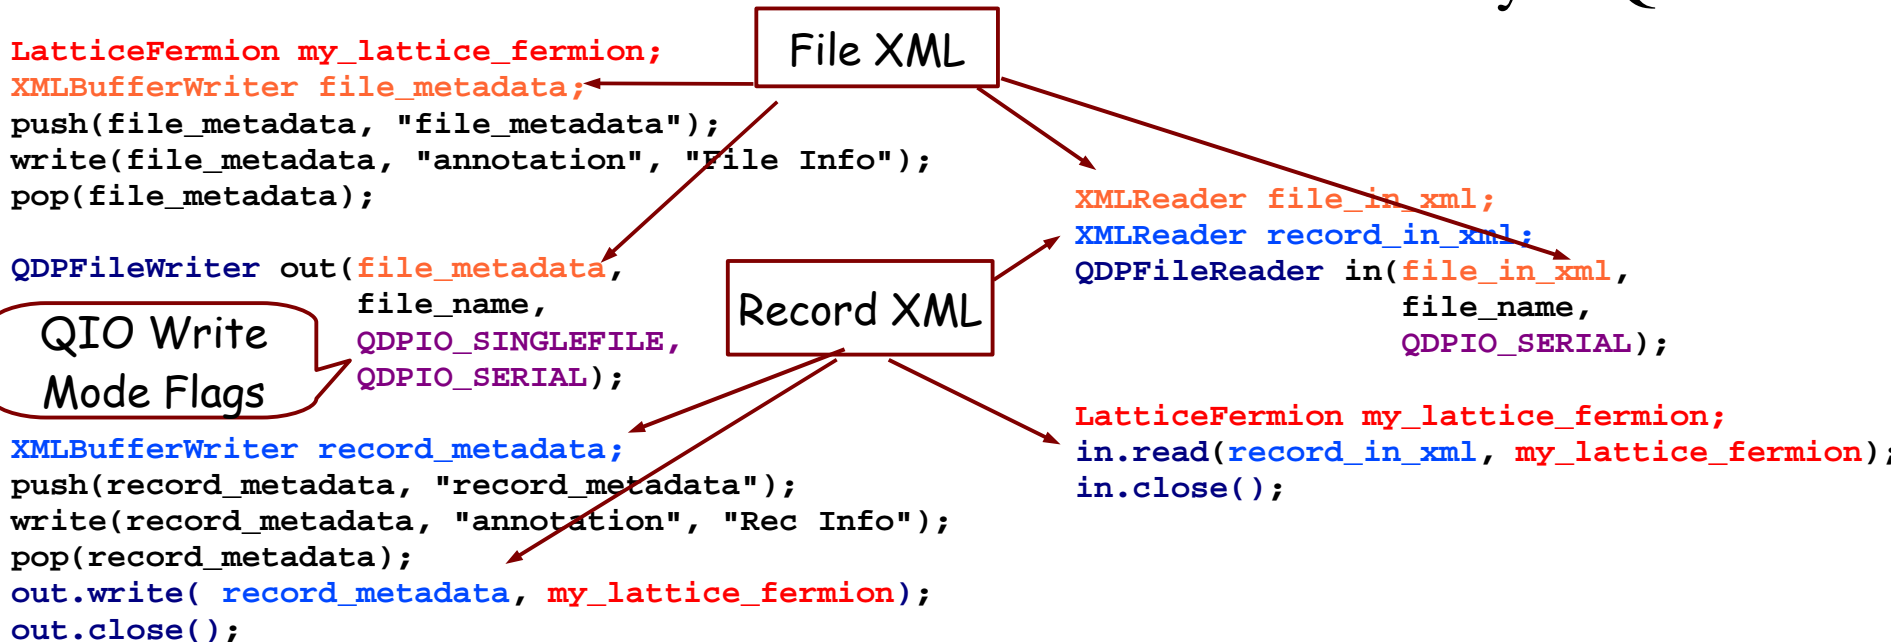


Writing/Reading  
patterns for  
QDPIO\_SERIAL



# QDP++ interface to QIO

- Write with QDPFileWriter
- Must supply user file and user record XML as XMLBufferWriter-s
- Read with QDPFileReader
- User File XML and User Record XML returned in XML Readers
- Checksum/ILDG details checked internally to QIO



# ILDG Support (?)

- The underlying QIO support layer can handle ILDG format.
- Specialization of write() function and type traits:
  - multi1d<LatticeColorMatrix> always written in ILDG format.
  - ILDG dataLFN is optional argument to the QDPFileWriter constructor.
- **nersc2ildg** program provided in examples/ directory
- **lime\_contents** and **lime\_extract\_record** programs from QIO automatically built and installed

# Custom Memory Allocation

- Occasionally need to allocate/free memory explicitly e.g. to provide memory to external library.
- Memory may need custom attributes (eg fast/communicable etc)
- Memory may need to be suitably aligned.
- May want to monitor memory usage

Allocate memory from desired pool if possible, with alignment suitable to pool

```
pointer=QDP::Allocator::theQDPAllocator::Instance().allocate( size,  
QDP::Allocator::FAST);
```

```
QDP::Allocator::theQDPAllocator::Instance()::free(pointer);
```

Namespace

Get reference to allocator

MemoryPoolHint (attribute)

# Move To Fast Memory (If you have it)

- QCDOC/SP inspired construct (copy to CRAM?)
- moves/copies data to/from fast memory (eg EDRAM)
- NOP on machines where there is no fast memory

```
LatticeFermion x;
```

```
moveToFastMemoryHint(x);
```

```
// Do some fast computation here
```

```
revertFromFastMemoryHint(x,true);
```

Accelerate x!  
Do not copy contents.  
Contents of x lost

Bring x back to slow  
memory. Copy contents

# Efficiency Considerations

- PETE eliminates Lattice sized temporaries
- But still there are temporaries at the site level
  - This really hurts performance
- Workaround: Cliché Expressions Optimized
  - eg:  $a*x + y$ ,  $a*x + P y$ ,  $\text{norm2}()$ ,  $\text{innerProduct}()$  etc.
  - Optimizations in *C*, SSE and bagel\_qdp library
- Non optimized expression still slow:
  - eg:  $a*x+y+z$  etc.

# Summary

- QDP++ and Chroma are quite mature now
- In this presentation I have discussed
  - basic usage (getting, building, testing, using)
  - a reasonably thorough description of QDP++
- Stay tuned for:
  - aspects of the chroma class structure
  - tutorial exercises
  - a presentation on writing XML input files